

1 Introduction

Currently, the security of networked infrastructures tends to be managed statically. That is, security requirements are developed and implemented within the networking environment, and all of the information that traverses the network will have these hard-coded security policies applied to it. Thus, the security policies applied to all information is "one-size-fits-all," and can typically only be modified by network administrators. The result is that if there is a mismatch between mission requirements and network security, the network resources simply cannot be used to service the mission.

The ability to quickly, securely, and reliably share information across security domains is a capability that must be developed in order to ensure mission success in Department of Defence (DoD) cloud-based network-centric environments. We propose the continued development of an assured information-sharing framework for cloud-based systems that leverages our ongoing work in the areas of policy-based usage management and semantic interoperability. The development of this framework involves research related to the creation of a novel approach to information sharing that treats security as a commodity that can be dynamically provisioned within the cloud, along with other cloud resources.

Policy Tools. First, we must develop the technologies that allow administrators to easily generate policies that leverage cloud computing-specific capabilities. That is, additional research is needed to develop a set of cloud computing-specific security policies that can be encoded into the extensible markup language (XML) policy format we currently use. This work will include extending our existing policy generator capabilities to address these cloud specific policies.

Dynamic Evaluation. In addition, the usage management framework must be extended to include the ability to formally interpret and enforce cloud-specific policies. These capabilities must include the ability to dynamically monitor the policies associated with resources in use in order to immediately react to changes. We envision this research including the development of a satellite usage management capability that operates in each VM. This approach will also require the development of capabilities that allow the global usage management mechanism to communicate in an assured manner with the usage management mechanisms executing in the individual VMs, providing comprehensive usage control across a distributed cloud environments.

Enforcement. Finally, policy enforcement mechanisms are needed at all levels of the cloud hierarchy to control how data and services are used. For example, if a data item is specified at a certain level of sensitivity, the usage management mechanisms must ensure that this item is not stored or routed through the cloud infrastructure in a way that would violate the policies associated with its sensitivity level. In order to accomplish this, our approach maintains clean separations between policy specification, interpretation and enforcement. This allows the framework to leverage security capabilities that are currently available or in development by other information security research groups.

Security Primitives and Control. In general, a DoD cloud computing infrastructure may be highly heterogeneous and used by multiple parties that may be involved in multiple missions, leading to highly complex usage patterns and security requirements. Thus, our Phase II research will continue to investigate the use of control theory in order to ensure that multiple missions can be simultaneously satisfied within a given infrastructure. Our next steps in this area involve investigating VM performance control capabilities that enable monitoring of central processing unit utilization and the development of appropriate control algorithms that automatically add or remove VMs to meet performance requirements within a given security posture.

Additional security issues associated with cloud computing that will be considered as a part of this usage management/control-theoretic framework include joint tenancy of a not trusted user's VM on the same server hardware, data remnants in cloud storage repositories and VMs, and the trust associated with cloud provider's administrators (who have complete access to user data and VMs). Some of these concerns could

potentially be addressed with usage policies; however, further research is needed to explore how the proposed usage management framework, in conjunction with

2 Operational Scenarios and Requirements

Nublu is intended to provide assured information sharing capabilities over private and public infrastructures. To do this, we will provide a secure virtualized system protected by integrated policy-based usage management integrated with key encryption technologies. Virtualization will span the entire system, from data repositories to communication networks to application servers to clients, when possible.

2.1 Abstract Requirements

We currently have a group of abstract requirements the system must fulfil to be useful. These address things like complete infrastructure virtualization to policy creation:

FiXme: What other abstract requirements do we want to list?

- **Flexible Infrastructure Provisioning** — Infrastructure must be able to be provisioned without human intervention.
- **Dynamic Infrastructure Provisioning** — Networks and systems must be able to be created and destroyed automatically based on need.
- **Data-driven Security Application** — Security services must be applied to information based on characteristics of that information, associated policies, and environmental context.
- **Automated Security Services** — Security services must be able to be applied without human intervention.
- **Automated Performance Tuning** — System performance must be tunable by the system without human intervention.
- **Virtualized Infrastructure** — Infrastructure must be completely virtualized in order to support dynamic allocation and provisioning.
- **Easy Policy Creation** — The system needs tools that enable automated policy creation whenever possible based on either none or some small amount of entered information.

2.2 Selected Scenarios

We have specific scenarios that embody information sharing that we can use to highlight the kind of functionality the system needs to produce in order to provide assured cloud-based information sharing. These scenarios have distributed communication requirements, use encryption, have multiple parties of various trust levels and need to know, and use both human and machine agents.

FiXme: These need to be elaborated, but how much?

Streaming Data From Drone. This is a data ingestion scenario. Data from some kind of drone, manned or unmanned, is pulled into a cloud-centric distribution system for storage and processing.

Streaming Data Consumption by Coalition Partner. In this scenario, a coalition partner is streaming video from a data source to some type of client. The video itself is not classified, but it is sensitive, and is accompanied by specific usage and routing policies.

Block Data Transfer within Datacenter. Datacenters in cloud environments usually host block storage for cloud clients on systems that are distinct from the computational systems themselves. Here, filesystem blocks are transferred from a block store to a client. The blocks themselves are categorized as sensitive, and have associated usage and management policies.

Automated Information Consumption by Remote System Agent. Here, a system agent on a remote system acquires data from a variety of cloud services for analysis. The agent itself is authenticated and authorized to access information, and the secondary data resulting from the analysis is saved and protected by usage policies that provide provenance information pointing to the original data used.

U.S. Operator Information Creation. In this scenario, a remote operator in the field creates some kind of data record. This record can be either an audio, video, or document-centric record. The information is saved with a default policy based on information gathered automatically from the creation location and user authentication information as well as a small amount of information entered by the user.

3 Solution Overview

We will use a simple model to begin development. We have three distinct layers — *application*, *logical*, and *physical* layers. This application layer consists of *data repositories*, *communication networks*, *application servers*, and *clients*. The logical layer contains *images* and *instances* running in *containers*, while the physical layer contains all physical hardware and software required to enable the logical layer.

These can be combined into common *patterns*, including:

- Intra-datacenter
- Inter-datacenter
- Datacenter to cache
- Datacenter to application server
- Datacenter to Client
- Cache to application server
- Cache to client

All these communication patterns have distinct *trust profiles*.

3.1 Major Features

The major system features differentiating this approach from other information sharing approaches is the scale of the system, the system's flexibility, and integrated content usage management.

System Scale. Unlike other information environments like SIPRNet or JWICS, this system will prototype the ability to store information from a single system at a variety of classification and sensitivity levels via integrated encryption. Furthermore, this system will operate as a cloud computing environment with the scalability that implies¹.

¹Peter Mell and Timothy Grance (Sept. 2011). *NIST SP 800-145: The NIST Definition of Cloud Computing*.

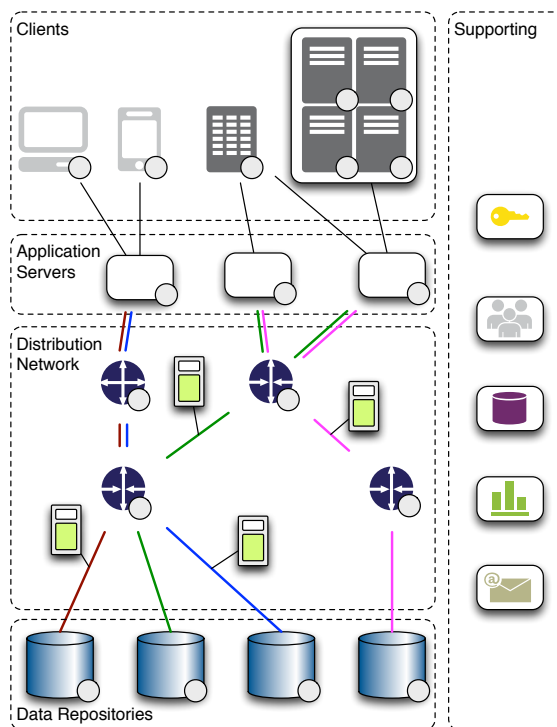


Figure 1: A notional view of the Nublu system, showing various required and provided services in an end-to-end perspective.

FiXme: Should this have a conops perspective? What's the appropriate depth?

System Flexibility. The system will allow policies to be evaluated within a dynamic context. For example, under cases of duress, the system may elect to transmit sensitive information to users that would generally not be allowed access to that information. This information would be shared in such a way that it could be retracted at a later date, and the system would create traceable records outlining which users received the information. In addition to this information management flexibility, the system will also support the kind of operational flexibility expected of cloud systems².

Content Usage Management. Nublu will enable management of information content throughout the system and information life-cycle.

References

Mell, Peter and Timothy Grance (Sept. 2011). *NIST SP 800-145: The NIST Definition of Cloud Computing*.

²Peter Mell and Timothy Grance (Sept. 2011). *NIST SP 800-145: The NIST Definition of Cloud Computing*.